

yugvegk33

February 17, 2026

Task 05: Personal Loan Acceptance Prediction.

Objective: Predict which customers are likely to accept a personal loan offer.

Dataset: Bank Marketing Dataset (UCI Machine learning Repository). About Dataset: UCI-Bank-Marketing-Dataset. The following information is drawn from the UCI Machine Learning Repository: <https://archive.ics.uci.edu/ml/datasets/bank+marketing>. Abstract: The data is related with direct marketing campaigns (phone calls) of a Portuguese banking institution. The classification goal is to predict if the client will subscribe a term deposit (variable y). Source: [Moro et al., 2014] S. Moro, P. Cortez and P. Rita. A Data-Driven Approach to Predict the Success of Bank Telemarketing. Decision Support Systems, Elsevier, 62:22-31, June 2014. Data Set Information: The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed. There are four datasets: Bank-additional-full.csv with all examples (41188) and 20 inputs, ordered by date (from May 2008 to November 2010), very close to the data analyzed in [Moro et al., 2014] <- this is used 2) bank-additional.csv with 10% of the examples (4119), randomly selected from 1), and 20 inputs. 3) bank-full.csv with all examples and 17 inputs, ordered by date (older version of this dataset with less inputs). 4) bank.csv with 10% of the examples and 17 inputs, randomly selected from 3 (older version of this dataset with less inputs). The classification goal is to predict if the client will subscribe (yes/no) a term deposit (variable y).

Load the Dataset:

```
[1]: import pandas as pd  
  
[2]: Dataset = pd.read_csv("bank-additional-full.csv")  
  
[3]: Dataset = pd.read_csv("bank-additional-full.csv", sep=";")
```

Data Exploration:

```
[4]: Dataset.head()  
  
[4]:   age      job marital education default housing loan    contact \
 0   56  housemaid  married  basic.4y      no      no  no  telephone
 1   57    services  married high.school  unknown      no  no  telephone
 2   37    services  married high.school      no     yes  no  telephone
 3   40    admin.  married  basic.6y      no      no  no  telephone
 4   56    services  married high.school      no      no  yes  telephone
```

```

month day_of_week ... campaign pdays previous poutcome emp.var.rate \
0 may mon ... 1 999 0 nonexistent 1.1
1 may mon ... 1 999 0 nonexistent 1.1
2 may mon ... 1 999 0 nonexistent 1.1
3 may mon ... 1 999 0 nonexistent 1.1
4 may mon ... 1 999 0 nonexistent 1.1

cons.price.idx cons.conf.idx euribor3m nr.employed y
0 93.994 -36.4 4.857 5191.0 no
1 93.994 -36.4 4.857 5191.0 no
2 93.994 -36.4 4.857 5191.0 no
3 93.994 -36.4 4.857 5191.0 no
4 93.994 -36.4 4.857 5191.0 no

```

[5 rows x 21 columns]

[5]: `Dataset.tail()`

```

[5]: age job marital education default housing loan \
41183 73 retired married professional.course no yes no
41184 46 blue-collar married professional.course no no no
41185 56 retired married university.degree no yes no
41186 44 technician married professional.course no no no
41187 74 retired married professional.course no yes no

contact month day_of_week ... campaign pdays previous \
41183 cellular nov fri ... 1 999 0
41184 cellular nov fri ... 1 999 0
41185 cellular nov fri ... 2 999 0
41186 cellular nov fri ... 1 999 0
41187 cellular nov fri ... 3 999 1

poutcome emp.var.rate cons.price.idx cons.conf.idx euribor3m \
41183 nonexistent -1.1 94.767 -50.8 1.028
41184 nonexistent -1.1 94.767 -50.8 1.028
41185 nonexistent -1.1 94.767 -50.8 1.028
41186 nonexistent -1.1 94.767 -50.8 1.028
41187 failure -1.1 94.767 -50.8 1.028

nr.employed y
41183 4963.6 yes
41184 4963.6 no
41185 4963.6 no
41186 4963.6 yes
41187 4963.6 no

```

```
[5 rows x 21 columns]
```

```
[6]: Dataset.describe()
```

```
[6]:
```

	age	duration	campaign	pdays	previous	\
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000	
mean	40.02406	258.285010	2.567593	962.475454	0.172963	
std	10.42125	259.279249	2.770014	186.910907	0.494901	
min	17.00000	0.000000	1.000000	0.000000	0.000000	
25%	32.00000	102.000000	1.000000	999.000000	0.000000	
50%	38.00000	180.000000	2.000000	999.000000	0.000000	
75%	47.00000	319.000000	3.000000	999.000000	0.000000	
max	98.00000	4918.000000	56.000000	999.000000	7.000000	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
count	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	0.081886	93.575664	-40.502600	3.621291	5167.035911
std	1.570960	0.578840	4.628198	1.734447	72.251528
min	-3.400000	92.201000	-50.800000	0.634000	4963.600000
25%	-1.800000	93.075000	-42.700000	1.344000	5099.100000
50%	1.100000	93.749000	-41.800000	4.857000	5191.000000
75%	1.400000	93.994000	-36.400000	4.961000	5228.100000
max	1.400000	94.767000	-26.900000	5.045000	5228.100000

```
[7]: Dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   age              41188 non-null   int64  
 1   job              41188 non-null   object  
 2   marital          41188 non-null   object  
 3   education        41188 non-null   object  
 4   default          41188 non-null   object  
 5   housing          41188 non-null   object  
 6   loan              41188 non-null   object  
 7   contact           41188 non-null   object  
 8   month             41188 non-null   object  
 9   day_of_week       41188 non-null   object  
 10  duration          41188 non-null   int64  
 11  campaign          41188 non-null   int64  
 12  pdays             41188 non-null   int64  
 13  previous          41188 non-null   int64  
 14  poutcome          41188 non-null   object  
 15  emp.var.rate      41188 non-null   float64
```

```
16 cons.price.idx 41188 non-null float64
17 cons.conf.idx 41188 non-null float64
18 euribor3m 41188 non-null float64
19 nr.employed 41188 non-null float64
20 y 41188 non-null object
dtypes: float64(5), int64(5), object(11)
memory usage: 6.6+ MB
```

```
[8]: Dataset.shape
```

```
[8]: (41188, 21)
```

```
[9]: Dataset.columns
```

```
[9]: Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
       'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
       'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
       'cons.conf.idx', 'euribor3m', 'nr.employed', 'y'],
      dtype='object')
```

```
[10]: Dataset.dtypes
```

```
[10]: age          int64
job           object
marital        object
education      object
default         object
housing         object
loan            object
contact         object
month           object
day_of_week     object
duration        int64
campaign        int64
pdays           int64
previous        int64
poutcome        object
emp.var.rate    float64
cons.price.idx float64
cons.conf.idx  float64
euribor3m      float64
nr.employed    float64
y               object
dtype: object
```

Perofrm basic data exploration on features such as age, job and martial status.

For job:

```
[11]: Dataset["job"].value_counts()
```

```
[11]: job
admin.          10422
blue-collar     9254
technician      6743
services         3969
management      2924
retired          1720
entrepreneur    1456
self-employed    1421
housemaid        1060
unemployed       1014
student           875
unknown           330
Name: count, dtype: int64
```

```
[12]: Dataset["job"].describe()
```

```
[12]: count      41188
unique       12
top          admin.
freq         10422
Name: job, dtype: object
```

For age:

```
[13]: Dataset["age"].value_counts()
```

```
[13]: age
31      1947
32      1846
33      1833
36      1780
35      1759
...
89      2
91      2
94      1
87      1
95      1
Name: count, Length: 78, dtype: int64
```

```
[14]: Dataset["age"].describe()
```

```
[14]: count      41188.00000
mean        40.02406
```

```
std           10.42125
min          17.00000
25%         32.00000
50%         38.00000
75%         47.00000
max          98.00000
Name: age, dtype: float64
```

For marital status:

```
[15]: Dataset["marital"].value_counts()
```

```
[15]: marital
      married    24928
      single     11568
      divorced   4612
      unknown     80
      Name: count, dtype: int64
```

```
[16]: Dataset["marital"].describe()
```

```
[16]: count      41188
      unique       4
      top        married
      freq      24928
      Name: marital, dtype: object
```

Basic Data Visualization:

```
[17]: import matplotlib as pyplot
from matplotlib import pyplot as plt
import seaborn as sns
```

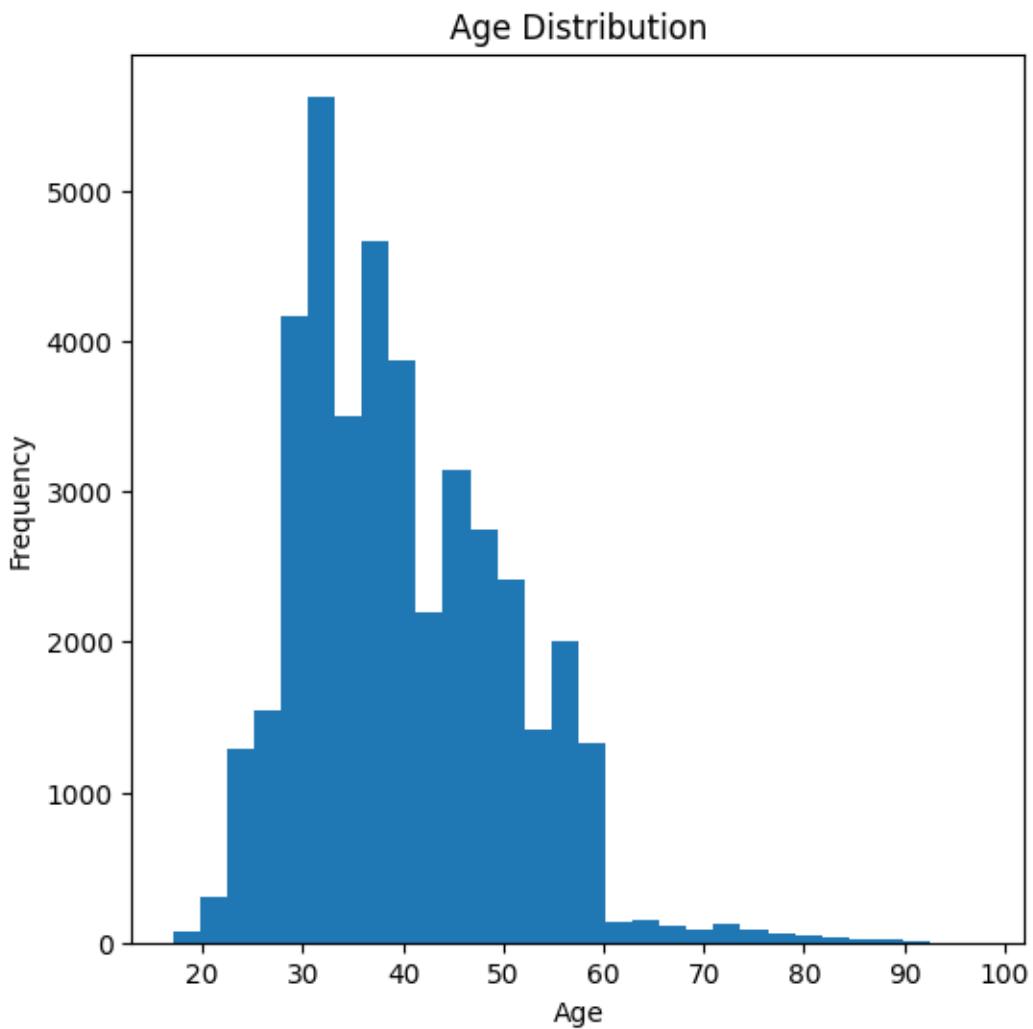
Age Histogram:

```
[18]: plt.figure(figsize = (6,6))

plt.hist(Dataset["age"], bins = 30)

plt.xlabel("Age")
plt.ylabel("Frequency")
plt.title("Age Distribution")

plt.show()
```



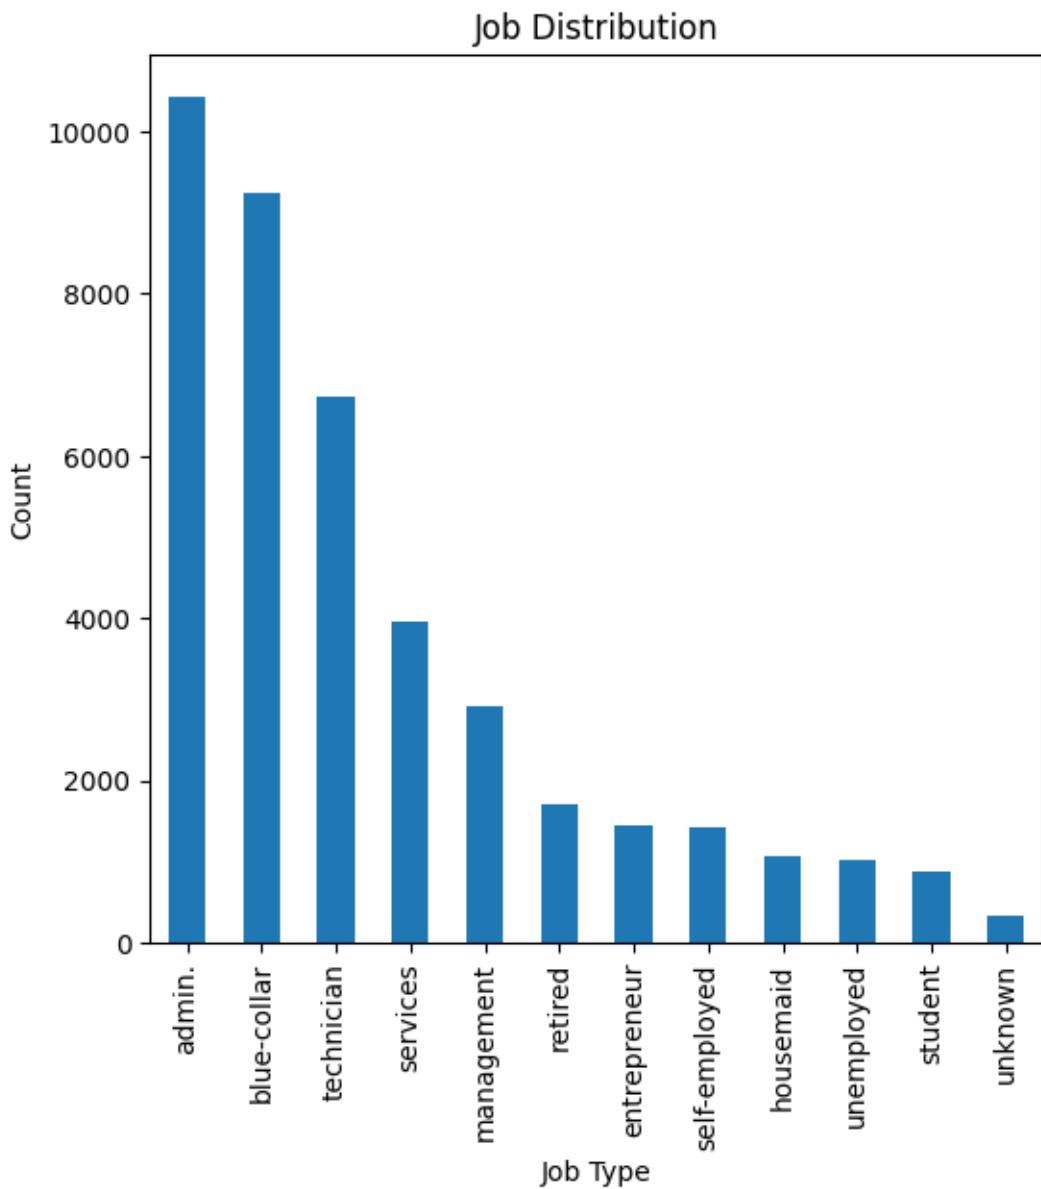
Job Count Plot:

```
[19]: plt.figure(figsize = (6,6))

Dataset["job"].value_counts().plot(kind="bar")

plt.xlabel("Job Type")
plt.ylabel("Count")
plt.title("Job Distribution")

plt.show()
```



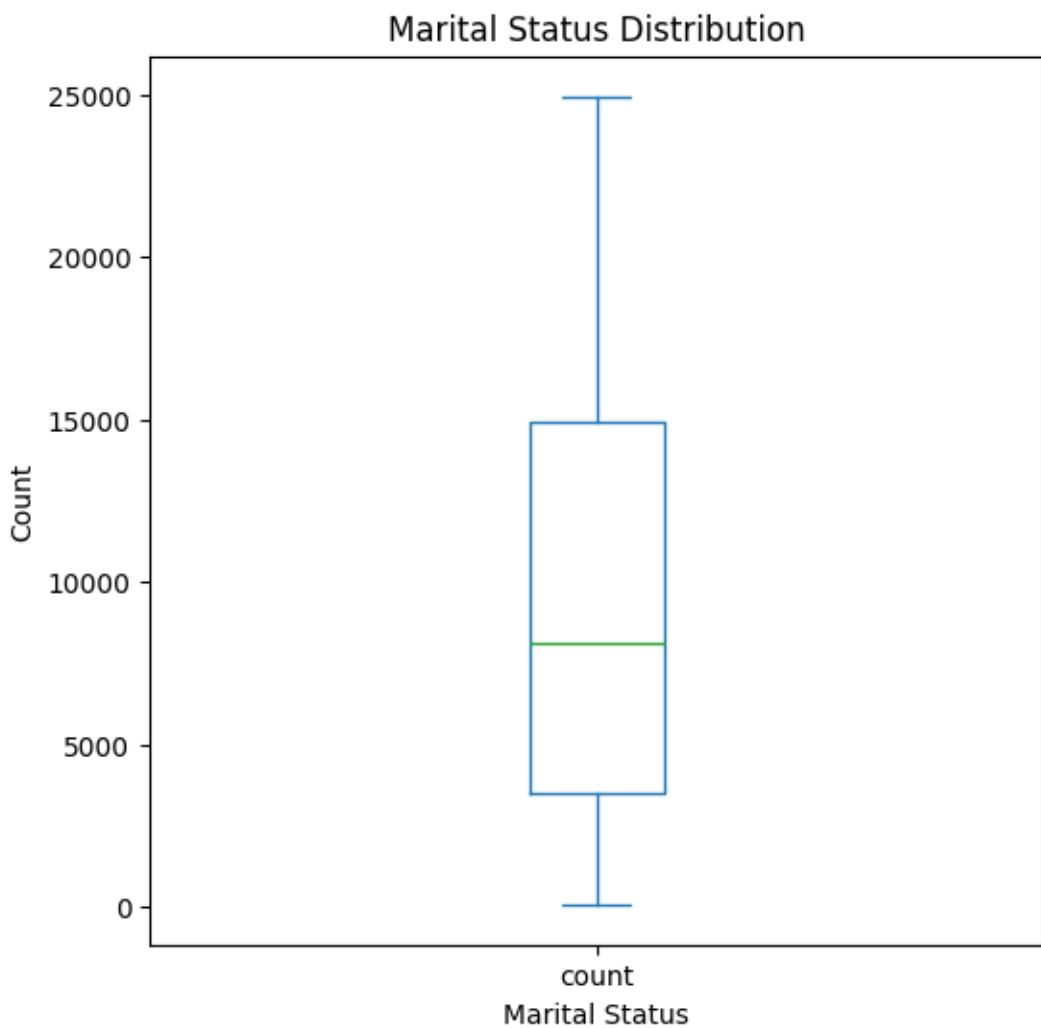
Marital Status Plot:

```
[20]: plt.figure(figsize = (6,6))

Dataset["marital"].value_counts().plot(kind="box")

plt.xlabel("Marital Status")
plt.ylabel("Count")
plt.title("Marital Status Distribution")

plt.show()
```



Pairplot:

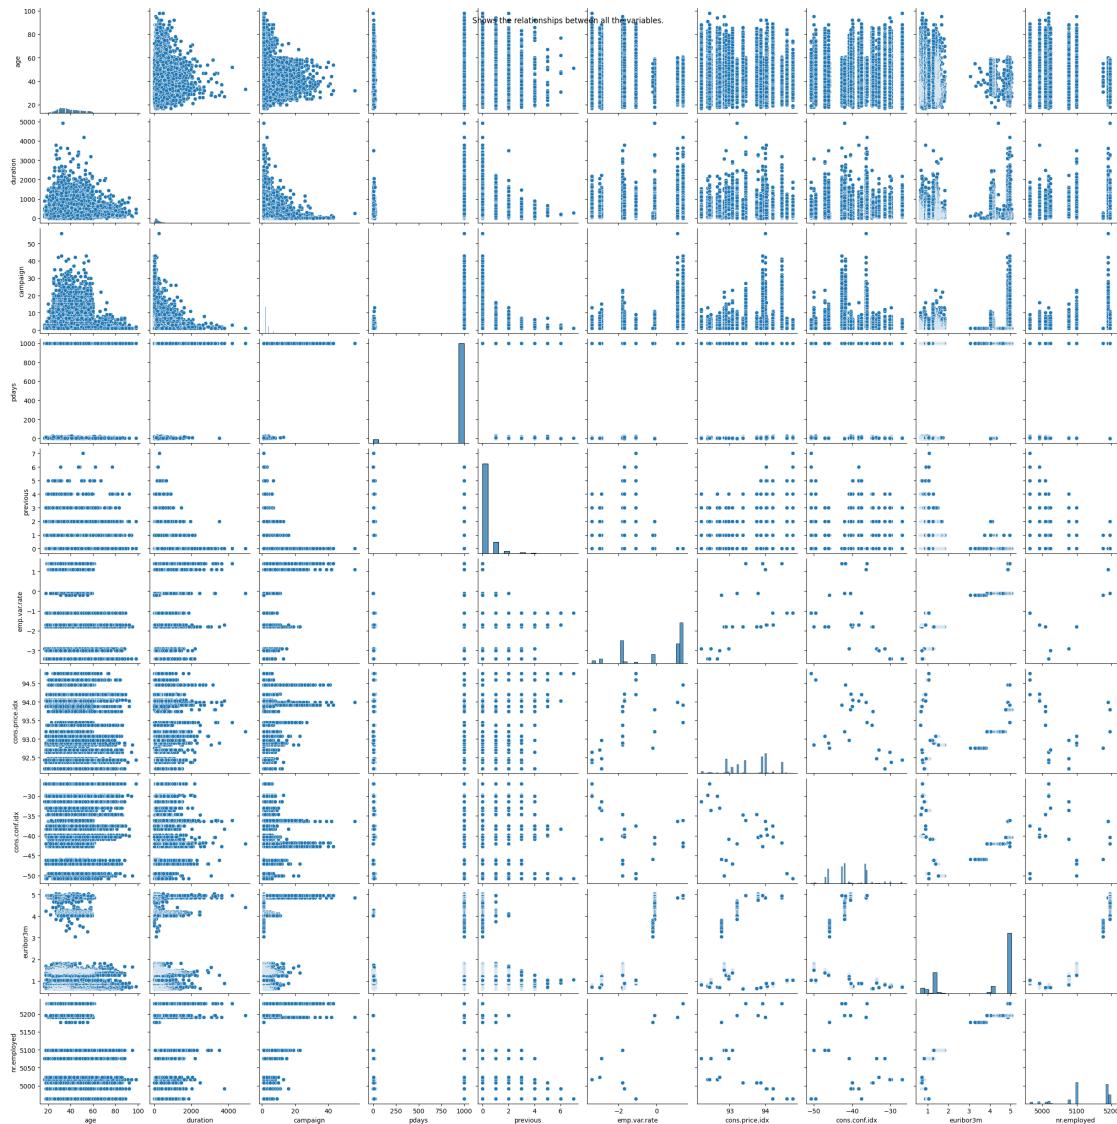
```
[21]: plt.figure(figsize=(6,6))

sns.pairplot(Dataset)

plt.suptitle("Shows the relationships between all the variables.")

plt.show()
```

<Figure size 600x600 with 0 Axes>



Analyze the results to identify which customer groups are likely to accept the offer.

Acceptance Rate by Job:

```
[22]: pd.crosstab(Dataset["job"], Dataset["y"], normalize="index") * 100
```

	y	no	yes
job			
admin.	87.027442	12.972558	
blue-collar	93.105684	6.894316	
entrepreneur	91.483516	8.516484	
housemaid	90.000000	10.000000	
management	88.782490	11.217510	
retired	74.767442	25.232558	

```

self-employed 89.514426 10.485574
services       91.861930  8.138070
student        68.571429 31.428571
technician     89.173958 10.826042
unemployed    85.798817 14.201183
unknown        88.787879 11.212121

```

Acceptance Rate by Marital Status:

```
[23]: pd.crosstab(Dataset["marital"], Dataset["y"], normalize="index") * 100
```

```

[23]: y          no         yes
      marital
divorced   89.679098 10.320902
married    89.842747 10.157253
single     85.995851 14.004149
unknown    85.000000 15.000000

```

Acceptance Rate by Age Groups:

```
[24]: Dataset["age_group"] = pd.cut(Dataset["age"], bins=[18,30,40,50,60,100])
```

```
pd.crosstab(Dataset["age_group"], Dataset["y"], normalize="index") * 100
```

```

[24]: y          no         yes
      age_group
(18, 30]   84.897959 15.102041
(30, 40]   90.253280  9.746720
(40, 50]   91.826172  8.173828
(50, 60]   89.346093 10.653907
(60, 100]  54.505495 45.494505

```

Handle Missing Values:

```
[25]: Dataset.isnull().sum()
```

```

[25]: age          0
      job          0
      marital      0
      education    0
      default      0
      housing      0
      loan          0
      contact      0
      month         0
      day_of_week   0
      duration      0
      campaign      0

```

```

pdays          0
previous       0
poutcome       0
emp.var.rate   0
cons.price.idx 0
cons.conf.idx  0
euribor3m      0
nr.employed    0
y              0
age_group      33
dtype: int64

```

[26]: Dataset = pd.get_dummies(Dataset, drop_first=True)

[27]: Dataset.head()

```

[27]:   age duration campaign pdays previous emp.var.rate cons.price.idx \
0     56      261         1    999        0        1.1      93.994
1     57      149         1    999        0        1.1      93.994
2     37      226         1    999        0        1.1      93.994
3     40      151         1    999        0        1.1      93.994
4     56      307         1    999        0        1.1      93.994

           cons.conf.idx euribor3m nr.employed ... day_of_week_thu \
0            -36.4      4.857      5191.0 ... False
1            -36.4      4.857      5191.0 ... False
2            -36.4      4.857      5191.0 ... False
3            -36.4      4.857      5191.0 ... False
4            -36.4      4.857      5191.0 ... False

           day_of_week_tue day_of_week_wed poutcome_nonexistent poutcome_success \
0             False           False                True            False
1             False           False                True            False
2             False           False                True            False
3             False           False                True            False
4             False           False               True            False

           y_yes age_group_(30, 40] age_group_(40, 50] age_group_(50, 60] \
0    False           False           False            True
1    False           False           False            True
2    False           True            False            False
3    False           True            False            False
4    False           False           False            True

           age_group_(60, 100]
0             False
1             False

```

```
2      False
3      False
4      False
```

```
[5 rows x 58 columns]
```

Train a Logistic Regression and Decision Tree Classifier.

Logistic Regression:

```
[28]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, confusion_matrix
```

```
[29]: x = Dataset.drop("y_yes", axis=1)
y = Dataset["y_yes"]
```

```
[30]: x = pd.get_dummies(x, drop_first=True)
```

```
[31]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,random_state=42)
```

```
[32]: print("Training Data:", x_train.shape)
print("Testing Data:", x_test.shape)
```

```
Training Data: (32950, 57)
Testing Data: (8238, 57)
```

```
[33]: logistic_model = LogisticRegression(max_iter=1000)
logistic_model.fit(x_train, y_train)

y_predict_logistic = logistic_model.predict(x_test)
```

```
c:\Users\zaigh\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result()
```

```
[34]: model = LogisticRegression()
model.fit(x_train, y_train)

c:\Users\zaigh\AppData\Local\Programs\Python\Python313\Lib\site-
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
    n_iter_i = _check_optimize_result()
```

[34]: LogisticRegression()

Evaluate the model using the accuracy and confusion matrix.

```
[35]: print("Logistic Regression Accuracy: \n",
accuracy_score(y_test, y_predict_logistic))

print("Confusion Matrix: \n",
confusion_matrix(y_test, y_predict_logistic))
```

Logistic Regression Accuracy:

0.9107793153678078

Confusion Matrix:

```
[[7106 197]
 [ 538 397]]
```

```
[36]: cm = confusion_matrix(y_test, y_predict_logistic)
TN, FP, FN, TP = cm.ravel()

# Create a Structured Table.
metrics_table = pd.DataFrame({
    'Metric': ['TP', 'TN', 'FP', 'FN', 'Accuracy', 'Precision', 'Recall'],
    'Value': [TP, TN, FP, FN,
              (TP+TN)/(TP+TN+FP+FN),
              TP/(TP+FP),
              TP/(TP+FN)]})

print(metrics_table)
```

	Metric	Value
0	TP	397.000000
1	TN	7106.000000
2	FP	197.000000

```

3      FN    538.000000
4  Accuracy    0.910779
5  Precision    0.668350
6  Recall     0.424599

```

Decision Tree:

```
[37]: decisionTree_model = DecisionTreeClassifier(random_state=42)

decisionTree_model.fit(x_train, y_train)

y_predict_decisionTree = decisionTree_model.predict(x_test)
```

Evaluate the model using the accuracy and confusion matrix.

```
[38]: print("Decision Tree Accuracy: \n",
           accuracy_score(y_test, y_predict_decisionTree))

print("Confusion Matrix: \n",
      confusion_matrix(y_test, y_predict_decisionTree))
```

```

Decision Tree Accuracy:
0.8869871327992231
Confusion Matrix:
[[6816 487]
 [ 444 491]]

```

```
[39]: cm = confusion_matrix(y_test, y_predict_decisionTree)
TN, FP, FN, TP = cm.ravel()

# Create a Structured Table.
metrics_table = pd.DataFrame({
    'Metric': ['TP', 'TN', 'FP', 'FN', 'Accuracy', 'Precision', 'Recall'],
    'Value': [TP, TN, FP, FN,
              (TP+TN)/(TP+TN+FP+FN),
              TP/(TP+FP),
              TP/(TP+FN)]})
print(metrics_table)
```

	Metric	Value
0	TP	491.000000
1	TN	6816.000000
2	FP	487.000000
3	FN	444.000000
4	Accuracy	0.886987
5	Precision	0.502045
6	Recall	0.525134

Skills: Data exploration and visualization. Classification modeling. Business insights extraction

from the data.

Task Completed. Best Wishes. Zaigham Abbas.